

Patternverse — Response Style Reference for Dev

Purpose: Side-by-side comparison of an actual live system response vs. a spec-compliant corrected version, annotated by which Master Build File section each issue/fix belongs to. Use this as the calibration reference when adjusting the runtime prompt's response-generation instructions.

Test input (unchanged, both columns respond to this):

"A colleague gave feedback on my project in a group meeting instead of privately. I felt a flash of embarrassment, then defensiveness, and I found myself over-explaining my reasoning for the next ten minutes even after they'd moved on. Later I noticed I was still replaying it that evening, drafting a mental rebuttal I never sent."

ACTUAL system response (live)

"It must have been tough to feel embarrassed and then have that escalate into defensiveness in front of everyone. Can you think of a specific recent moment when a colleague gave you feedback in a group setting, even if it wasn't in the same way as this time?"

CORRECTED response (reference)

mirror: A colleague gave feedback on your project in a group meeting rather than privately. You noticed embarrassment, then defensiveness, then continued explaining your reasoning for about ten minutes after the conversation had moved on. That evening, the moment replayed, and you drafted a response you didn't send.

rca_frame: Trigger: public feedback in a group setting. Reaction: embarrassment followed by extended justification, plus continued mental rehearsal afterward. Outcome: no resolution reached — the unsent message and evening replay suggest the loop stayed open rather than closing.

pattern_hypothesis: Not enough is established yet to classify a loop state. What's observable so far is a justification response that continued past the point where it was being received — outlasting the triggering moment itself.

context_note: This is based on a single instance. I don't yet know whether this response

shows up in other settings, or whether something similar has happened before in a different context.

directive: Has a moment like this come up before — maybe in a different setting, with a different person?

Annotated Differences — Mapped to Build File Sections

Issue in ACTUAL response	Build file section violated	What CORRECTED does instead	Build file section to enforce fix
"It must have been tough..." — unearned empathy/validation, no MSL-1 confirmed	Part C.6, CIM Class B6 (empathy-lead clause); Part C.6 Enforcement clause (3.5 amendment)	No validation language. Mirror section states only observable facts.	Part B.3, Runtime Prompt → mirror instruction. Add explicit rule: mirror section may not contain emotionally validating language ("must have been," "that's hard," "tough") unless <code>msl_primary = msl_1_trauma</code> confirmed.
"...have that escalate into defensiveness" — causal claim stated as settled fact	Part C.6, CIM Class B6 (restatement clause); also overlaps Part B.3 CIM check "B2 — Mirror section contains interpretive content rather than observable description"	<code>mirror</code> lists embarrassment and defensiveness as parallel observed facts, not a causal chain. <code>pattern_hypothesis</code> explicitly defers classification instead of asserting it.	Part B.3, Runtime Prompt → mirror and pattern_hypothesis instructions. Reinforce B2's existing rule and extend it to cover causal-link assertions, not just labels.
			Part B.3, Output Format section. Confirm the live

No five-section structure — single casual paragraph blending mirror, interpretation, and directive together	Part B.3, Output Format ("always include all five sections")	Five sections clearly separated: mirror / rca_frame / pattern_hypothesis / context_note / directive	runtime prompt is actually enforcing five distinct sections — the live app's casual tone suggests the front-end may be flattening or the prompt may not be structurally enforcing this. Cross-check against Part F.8 (Result Screen must display all five sections) — if sections aren't being generated distinctly, F.8 can't be implemented correctly either.
<code>pattern_hypothesis</code> skipped entirely — no acknowledgment that confidence is low / classification isn't earned	Part B.3, Confidence Gating ("fewer than 2 answered dimensions → reflect or qualify only. No classification.")	<code>context_note</code> explicitly states only one instance is known and classification is deferred	Part B.3, Runtime Prompt → step 5 (confidence gating) and step 8 (output formatting). The corrected response makes the gating state visible to the user via <code>context_note</code> ; the live response gives no signal that the system is still in low-confidence territory.
Directive question is asked correctly (second-instance probe) — this part is not broken	Part B.2, Follow-Up Question Selection Logic	Same intent, same dimension target, phrased as a plain question without wrapping it in a restated interpretation	No fix needed here — this confirms the question-selection logic itself works; the defect is isolated to response phrasing/structure, not to dimension targeting.

Net instruction for dev

The defect is concentrated entirely in **Part B.3's Runtime Prompt**, specifically the instructions governing the `mirror` and `pattern_hypothesis` sections, plus whatever is causing the five-section structure to collapse into one casual paragraph in the live app. The Follow-Up Question Selection Logic (Part B.2) and CIM Class B5 (Part C.6) are both confirmed working correctly elsewhere in testing — leave those untouched. This is a targeted prompt-language fix, not a pipeline or architecture change.